


# 网络实验报告

2016K8009909005 雷正宇

## 广播网络实验


### 实验内容

实现节点广播的broadcast\_packet函数；




验证广播网络能够正常运行：从一个端节点ping另一个端节点；

验证广播网络的效率：在three\_nodes\_bw.py进行iperf测量，两种场景：H1: iperf client; H2, H3: servers；H1: iperf server; H2, H3: clients；



先构建环形拓扑，验证该拓扑下节点广播会产生数据包环路。



## 实验流程

### 1. 实现节点广播函数

```

1 void broadcast_packet(iface_info_t *iface, const char *packet, int len)
2 {
3     iface_info_t *pos = NULL;
4     list_for_each_entry(pos, &instance->iface_list, list) {
5         if (pos != iface)
6             iface_send_packet(pos, packet, len);
7     }
8
9     fprintf(stdout, "TODO: broadcast packet here.\n");
10 }

```

这个函数用到了列表的遍历宏：

```

1 #define list_for_each_entry(pos, head, member) \
2     for (pos = list_entry((head)->next, typeof(*pos), member); \
3         &pos->member != (head); \
4         pos = list_entry(pos->member.next, typeof(*pos), member))

```

这里的 list 结构是 Linux 风格的：如果某个结构体类型需要当作被链表处理，就让它包含一个链表节点结构，从而被一个链表串起来。list\_entry 宏的作用是通过链表节点找到持有它的外部结构。如果不对结点进行删除操作，用 list\_for\_each\_entry 这个宏就能很方便地遍历整个链表。

编写完成后，利用 Makefile 编译即可完成 hub 的制作。

## 2. 验证广播网络

先将 three\_nodes\_bw.py 文件中的 hub 节点后台程序换成刚刚编写的版本：

```

1 # three_nodes_bw.py
2 # ...
3 # b1.cmd('./hub-reference &')
4 b1.cmd('./hub &')

```

然后启动 mininet，看看 h1, h2, h3 能不能相互 ping 通：

```
"Node: h1"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.2 -c 4
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.316 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.125 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.155 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.173 ms

--- 10.0.0.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3069ms
rtt min/avg/max/mdev = 0.125/0.192/0.316/0.074 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.3 -c 4
PING 10.0.0.3 (10.0.0.3) 56(84) bytes of data.
64 bytes from 10.0.0.3: icmp_seq=1 ttl=64 time=0.368 ms
64 bytes from 10.0.0.3: icmp_seq=2 ttl=64 time=0.171 ms
64 bytes from 10.0.0.3: icmp_seq=3 ttl=64 time=0.166 ms
64 bytes from 10.0.0.3: icmp_seq=4 ttl=64 time=0.376 ms

--- 10.0.0.3 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3058ms
rtt min/avg/max/mdev = 0.166/0.270/0.376/0.102 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

```
"Node: h2"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.1 -c 4
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=0.324 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=0.133 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=0.329 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=64 time=0.169 ms

--- 10.0.0.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3049ms
rtt min/avg/max/mdev = 0.133/0.238/0.329/0.090 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.3 -c 4
PING 10.0.0.3 (10.0.0.3) 56(84) bytes of data.
64 bytes from 10.0.0.3: icmp_seq=1 ttl=64 time=0.118 ms
64 bytes from 10.0.0.3: icmp_seq=2 ttl=64 time=0.140 ms
64 bytes from 10.0.0.3: icmp_seq=3 ttl=64 time=0.206 ms
64 bytes from 10.0.0.3: icmp_seq=4 ttl=64 time=0.174 ms

--- 10.0.0.3 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3073ms
rtt min/avg/max/mdev = 0.118/0.159/0.206/0.035 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

```
"Node: h3"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.1 -c 4
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
64 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=0.817 ms
64 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=0.127 ms
64 bytes from 10.0.0.1: icmp_seq=3 ttl=64 time=0.142 ms
64 bytes from 10.0.0.1: icmp_seq=4 ttl=64 time=0.165 ms

--- 10.0.0.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3033ms
rtt min/avg/max/mdev = 0.127/0.312/0.817/0.292 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# ping 10.0.0.2 -c 4
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.324 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.498 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.231 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.407 ms

--- 10.0.0.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3056ms
rtt min/avg/max/mdev = 0.231/0.365/0.498/0.098 ms
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

结果显示相互 ping 的过程是正常的。

### 3. 验证广播网络的效率

#### h2, h3 作为服务器

命令如下：

```
1 | h2 # iperf -s
2 | h3 # iperf -s
```

h1 连通 h2 和 h3:

```
"Node: h1"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# iperf -c 10.0.0.2 -t 30
connect failed: Connection refused
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# iperf -c 10.0.0.2 -t 30
-----
Client connecting to 10.0.0.2, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.1 port 41398 connected with 10.0.0.2 port 5001
[ ID] Interval          Transfer      Bandwidth
[ 13] 0.0-30.6 sec    655 KBytes   175 Kbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

```
"Node: h1"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast# iperf -c 10.0.0.3 -t 30
-----
Client connecting to 10.0.0.3, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.1 port 46554 connected with 10.0.0.3 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.7 sec   658 KBytes  175 Kbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast#
```

## h1 作为服务器

```
1 | h1 # iperf -s
```


h2, h3 连 h1:

```
"Node: h2"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast# iperf -c 10.0.0.1 -t 30
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.2 port 56156 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.8 sec   641 KBytes  170 Kbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast#
```

```
"Node: h3"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast# iperf -c 10.0.0.1 -t 30
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.3 port 38136 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.5 sec   641 KBytes  172 Kbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast#
```

## 4. 环路转发

在上面这个 naive 广播网络中，如果出现环路，就会无限转发。以如下网络拓扑为例：



构建该网络的代码如下：

circle.py

```
1  #!/usr/bin/python
2
3  import sys
4  import os.path
5
6  from mininet.topo import Topo
7  from mininet.net import Mininet
8  from mininet.link import TCLink
9  from mininet.cli import CLI
10
11 # Mininet will assign an IP address for each interface of a node
12 # automatically, but hub or switch does not need IP address.
13 def clearIP(n):
14     for iface in n.intfList():
15         n.cmd('ifconfig %s 0.0.0.0' % (iface))
16
17 class BroadcastTopo(Topo):
18     def build(self):
19         h1 = self.addHost('h1')
20         h2 = self.addHost('h2')
21
22         b1 = self.addHost('b1')
23         b2 = self.addHost('b2')
24         b3 = self.addHost('b3')
25
26         self.addLink(h1, b1, bw=10)
27         self.addLink(h2, b2, bw=10)
```

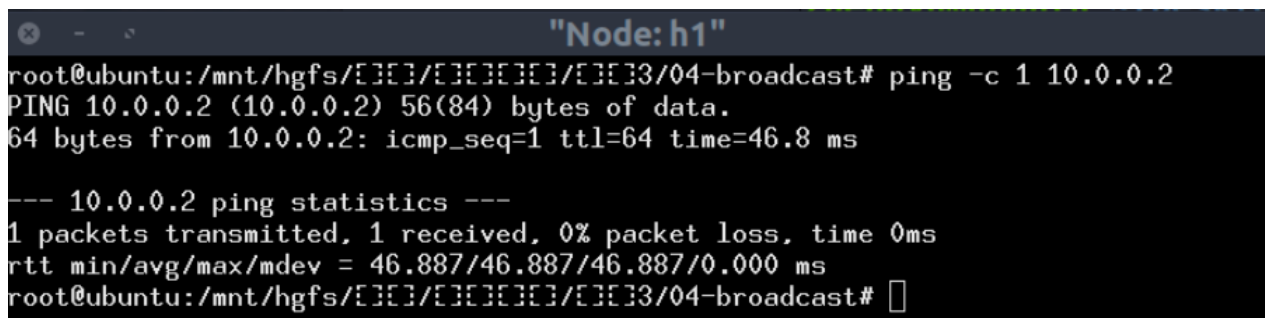
```

28         self.addLink(b1, b2, bw=10)
29         self.addLink(b1, b3, bw=10)
30         self.addLink(b2, b3, bw=10)
31
32     if __name__ == '__main__':
33         if not os.path.exists('/sbin/ethtool'):
34             print('ethtool not found, please install it using `apt install
ethtool`')
35             sys.exit(1)
36
37         topo = BroadcastTopo()
38         net = Mininet(topo = topo, link = TCLink, controller = None)
39
40         h1, h2, b1, b2, b3 = net.get('h1', 'h2', 'b1', 'b2', 'b3')
41         h1.cmd('ifconfig h1-eth0 10.0.0.1/8')
42         h2.cmd('ifconfig h2-eth0 10.0.0.2/8')
43
44         clearIP(b1)
45         clearIP(b2)
46         clearIP(b3)
47
48         for h in [h1, h2]:
49             h.cmd('./disable_offloading.sh')
50             h.cmd('./disable_ipv6.sh')
51
52         net.start()
53
54         b1.cmd('./hub &')
55         b2.cmd('./hub &')
56         b3.cmd('./hub &')
57         # b1.cmd('./hub-reference &')
58         # b2.cmd('./hub-reference &')
59         # b3.cmd('./hub-reference &')
60         CLI(net)
61         net.stop()

```

(其中大部分内容 copy 自 three\_nodes\_bw.py)

现在在 h1 节点 ping h2:



```

"Node: h1"
root@ubuntu:/mnt/hgfs/[[[[]/[[[[][[[[]/[[[[]3/04-broadcast# ping -c 1 10.0.0.2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=46.8 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 46.887/46.887/46.887/0.000 ms
root@ubuntu:/mnt/hgfs/[[[[]/[[[[][[[[]/[[[[]3/04-broadcast# 

```

在 h2 打开 wireshark 跟踪网络包, 会发现 h2 不断地在接收网络包, 截取了其中一个部分如下:



18874	1.006950534	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18875	1.006951125	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18876	1.006998250	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18877	1.007030253	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18878	1.007057307	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18879	1.007093121	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18880	1.007442370	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18881	1.007442978	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18882	1.007443607	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18883	1.007444085	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18884	1.007444556	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18885	1.007445263	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18886	1.007445767	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18887	1.007446256	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18888	1.007446801	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18889	1.007447505	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18890	1.007937953	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89
18891	1.007939475	f6:ba:77:b5:c6:89	8a:1f:66:25:77:d2	ARP	42	10.0.0.2	is	at	f6:ba:77:b5:c6:89

## 实验结果与分析

本次实验主要还是在于过程体验，需要编写代码、构建网络的部分并不是很多，所以很大部分实验结果已经在流程中表达出来了。

这次实验中遇到了一个当时（这部分报告撰写之前）发现的问题，一个是广播网络效率问题：我的网络广播 iperf 测试的带宽很低。后来我将节点广播函数中的 fprintf 函数的调用删去以后：

```

1 void broadcast_packet(iface_info_t *iface, const char *packet, int len)
2 {
3     iface_info_t *pos = NULL;
4     list_for_each_entry(pos, &instance->iface_list, list) {
5         if (pos != iface)
6             iface_send_packet(pos, packet, len);
7     }
8
9     // fprintf(stdout, "TODO: broadcast packet here.\n");
10 }

```

用修改后的 hub 完成的实验呈现的结果终于与 reference 相符：

```

root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/04-broadcast# iperf -c 10.0.0.2 -t 30
-----
Client connecting to 10.0.0.2, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.1 port 40086 connected with 10.0.0.2 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13]  0.0-30.1 sec  33.8 MBytes  9.40 Mbits/sec

```

这是 h1 作为 client 连接 h2 的带宽：9.40Mbits/s. h1 连接 h3 server 带宽与此值相近。以下是 h1 作为服务器的情形：

```
"Node: h2"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# iperf -c 10.0.0.1 -t 30
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.2 port 58762 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.1 sec  33.9 MBytes  9.43 Mbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

和 h2 作为服务器几乎没有区别。当 h2, h3 同时连接 h1 时:

```
"Node: h3"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# iperf -c 10.0.0.1 -t 30
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.3 port 55080 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.2 sec  32.1 MBytes  8.93 Mbits/sec
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast#
```

```
"Node: h2"
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/3/04-broadcast# iperf -c 10.0.0.1 -t 30
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 13] local 10.0.0.2 port 58766 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 13] 0.0-30.1 sec  33.8 MBytes  9.41 Mbits/sec
```

与 h1 作为服务器的情形几乎也相同, h3 先于 h2 执行 iperf, 最终带宽稍低一些, 总体上都小于 10Mbits/s。