

网络实验报告

雷正宇 leizhengyu16@mails.ucas.ac.cn

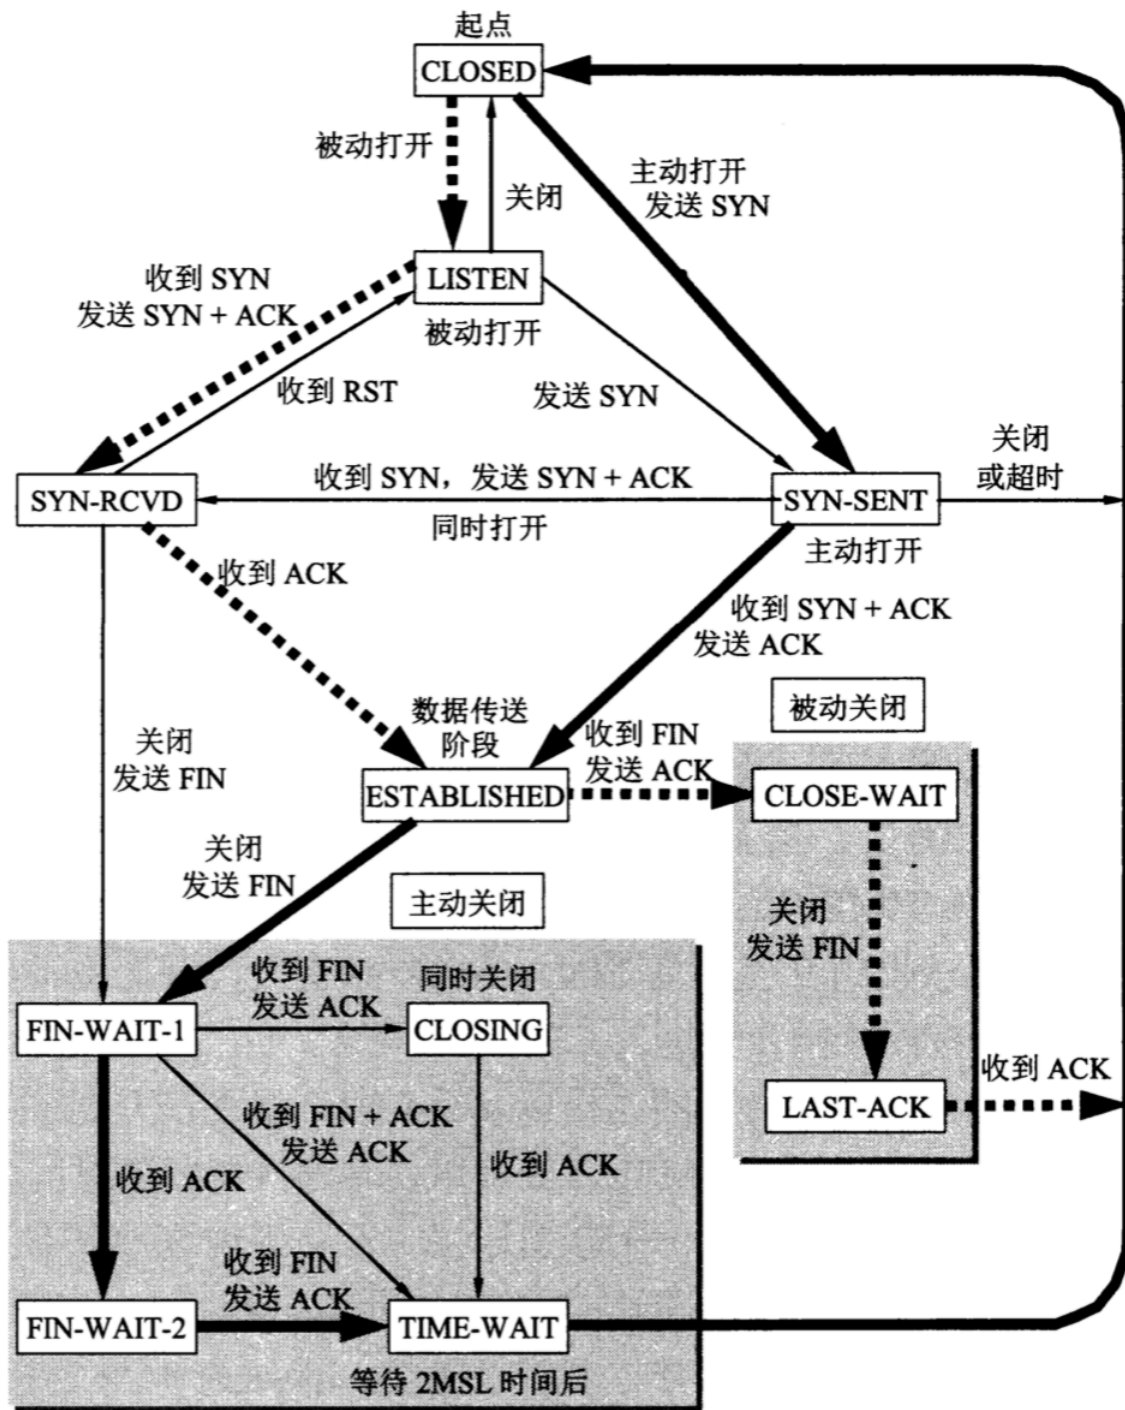
实验内容

实现 TCP 管理相关函数：

```
struct tcp_sock *alloc_tcp_sock();  
  
int tcp_sock_bind(struct tcp_sock *, struct sock_addr *);  
  
int tcp_sock_listen(struct tcp_sock *, int);  
  
int tcp_sock_connect(struct tcp_sock *, struct sock_addr *);  
  
struct tcp_sock *tcp_sock_accept(struct tcp_sock *);  
  
void tcp_sock_close(struct tcp_sock *);
```

实验流程

本次试验要完成一个具备建立链接和断开链接功能的 TCP 协议栈。首先要明确 TCP 状态机的运作方式：



其中每一个箭头过程都需要处理。

处理控制信号过程

```

1 void tcp_process(struct tcp_sock *tsk, struct tcp_cb *cb, char *packet)
2 {
3     // fprintf(stdout, "TODO: implement %s please.\n", __FUNCTION__);
4     u8 flags = cb->flags;
5
6     if (flags & TCP_RST) {
7         tcp_sock_close(tsk);
8         free_tcp_sock(tsk);
9         return;

```

```

10     }
11
12     log(DEBUG, IP_FMT ":%hu current state is %s.", \
13         HOST_IP_FMT_STR(tsk->sk_sip), tsk->sk_sport, \
14         tcp_state_str[tsk->state]);
15
16     switch (tsk->state) {
17     case TCP_CLOSED:
18         tcp_send_reset(cb);
19         return;
20     case TCP_LISTEN:
21         if (flags & TCP_SYN) {
22             struct tcp_sock *new = alloc_tcp_sock();
23             new->parent = tsk;
24             new->sk_dip = cb->saddr;
25             new->sk_dport = cb->sport;
26             new->sk_sip = cb->daddr;
27             new->sk_sport = cb->dport;
28
29             new->iss = tcp_new_iss();
30             new->snd_una = new->iss;
31             new->snd_nxt = new->iss;
32             new->rcv_nxt = cb->seq + 1;
33
34             tcp_set_state(new, TCP_SYN_RECV);
35
36             tcp_hash(new);
37             list_add_head(&new->list, &tsk->listen_queue);
38
39             tcp_send_control_packet(new, TCP_ACK | TCP_SYN);
40         } else {
41             log(ERROR, "invalid tcp packet while in state LISTEN");
42         }
43         return;
44     case TCP_SYN_RECV:
45         if (flags & TCP_ACK) {
46             list_delete_entry(&tsk->list);
47             tcp_sock_accept_enqueue(tsk);
48             wake_up(tsk->parent->wait_accept);
49             tcp_set_state(tsk, TCP_ESTABLISHED);
50         } else {
51             log(ERROR, "invalid tcp packet while in state SYN_RECV");
52         }
53         return;
54     case TCP_SYN_SENT:
55         if (flags & (TCP_SYN | TCP_ACK)) {
56             tsk->rcv_nxt = cb->seq + 1;
57             tsk->snd_nxt = cb->ack;
58

```

```

59     tcp_send_control_packet(tsk, TCP_ACK);
60     tcp_set_state(tsk, TCP_ESTABLISHED);
61     wake_up(tsk->wait_connect);
62 } else if (flags & TCP_SYN) {
63     tcp_send_control_packet(tsk, TCP_SYN | TCP_ACK);
64     tcp_set_state(tsk, TCP_SYN_RECV);
65 } else {
66     log(ERROR, "invalid tcp packet while in state SYN_SENT");
67 }
68 return;
69 case TCP_ESTABLISHED:
70     if (flags & TCP_FIN) {
71         tsk->rcv_nxt = cb->seq_end;
72         tcp_send_control_packet(tsk, TCP_ACK);
73         tcp_set_state(tsk, TCP_CLOSE_WAIT);
74     } else {
75         log(ERROR, "invalid tcp packet while in state ESTABLISHED");
76     }
77     return;
78 case TCP_CLOSE_WAIT:
79     log(ERROR, "invalid tcp packet while in state CLOSE_WAIT");
80     return;
81 case TCP_LAST_ACK:
82     if (flags & TCP_ACK) {
83         tcp_set_state(tsk, TCP_CLOSED);
84         tcp_unhash(tsk);
85     } else {
86         log(ERROR, "invalid tcp packet while in state LAST_ACK");
87     }
88     return;
89 case TCP_FIN_WAIT_1:
90     if (flags & TCP_ACK) {
91         tcp_set_state(tsk, TCP_FIN_WAIT_2);
92     } else if (flags & TCP_FIN) {
93         tcp_send_control_packet(tsk, TCP_ACK);
94         tcp_set_state(tsk, TCP_CLOSING);
95     } else {
96         log(ERROR, "invalid tcp packet while in state FIN_WAIT_1");
97     }
98     return;
99 case TCP_FIN_WAIT_2:
100     if (flags & TCP_FIN) {
101         tsk->rcv_nxt = cb->seq_end;
102         tcp_send_control_packet(tsk, TCP_ACK);
103         tcp_set_state(tsk, TCP_TIME_WAIT);
104         tcp_set_timewait_timer(tsk);
105     } else {
106         log(ERROR, "invalid tcp packet while in state FIN_WAIT_2");
107     }

```

```

108     case TCP_CLOSING:
109         if (flags & TCP_ACK) {
110             tcp_set_state(tsk, TCP_TIME_WAIT);
111             tcp_set_timewait_timer(tsk);
112         } else {
113             log(ERROR, "invalid tcp packet while in state FIN_WAIT_2");
114         }
115         return;
116     case TCP_TIME_WAIT:
117         log(ERROR, "invalid tcp packet while in state TIME_WAIT");
118         return;
119     default:
120         break;
121 }
122 }

```

该函数篇幅很长，但逻辑很清晰——先判断 TCP 目前的状态，再根据收到的控制信号进行相应的动作。其中部分动作只是改变此时的状态和发送对应控制信号，而少数动作则需要复杂操作。

当服务器进入 LISTEN 状态后，如果收到 SYN 信号，则需要新建一个子套接字，将其状态置为 SYN_RECV 并加入自己的监听队列。对于已经处于 SYN_RECV 状态的套接字，收到 ACK 包后要被唤醒。关于睡眠与唤醒的内容，下面还有叙述。接下来的几个过程都将和该过程有关联。

主动建立连接过程

```

1  int tcp_sock_connect(struct tcp_sock *tsk, struct sock_addr *skaddr)
2  {
3      // fprintf(stdout, "TODO: implement %s please.\n", __FUNCTION__);
4      tsk->sk_dip = sock_addr_get_ip(skaddr);
5      tsk->sk_dport = sock_addr_get_port(skaddr);
6      iface_info_t *iface = list_entry(instance->iface_list.next,
7      iface_info_t, list);
8      tsk->sk_sip = iface->ip;
9      tsk->sk_sport = tcp_get_port();
10
11      tcp_bind_hash(tsk);
12
13      tcp_send_control_packet(tsk, TCP_SYN);
14      tcp_set_state(tsk, TCP_SYN_SENT);
15      tcp_hash(tsk); // sequence?
16
17      sleep_on(tsk->wait_connect);
18
19      return 0;
20 }

```

过程分为4步：设置源/目的-端口/地址四元组，绑定到 bind_table，发送控制信号并切换状态，阻塞（睡眠）到 wait_connect 队列上。其中注意网络-本地字节序转换。等到收到 SYN + ACK 信号后，连接建立，该套接字进入 ESTABLISHED 状态。

被动监听建立过程

```
1 int tcp_sock_listen(struct tcp_sock *tsk, int backlog)
2 {
3     // fprintf(stdout, "TODO: implement %s please.\n", __FUNCTION__);
4     tsk->backlog = backlog;
5     tcp_set_state(tsk, TCP_LISTEN);
6     return tcp_hash(tsk);
7 }
```

该过程也很简单，修改套接字状态并将套接字塞入哈希表即可。

接收过程

```
1 struct tcp_sock *tcp_sock_accept(struct tcp_sock *tsk)
2 {
3     // fprintf(stdout, "TODO: implement %s please.\n", __FUNCTION__);
4     while (list_empty(&tsk->accept_queue)) {
5         sleep_on(tsk->wait_accept);
6     }
7
8     return tcp_sock_accept_dequeue(tsk);
9 }
```

一个被动建立连接方处理过程是这样的：从主套接字的 `accept_queue` 中获取队列头的子套接字，然后处理该子套接字。一般来说，以 server 为例（本实验不是这样的例子），一个 server 在建立主套接字后会不断针对主套接字的 `accept_queue` 进行处理，而封装调用即是 `tcp_sock_accept`。回到状态机处理信号的过程中，在本实验中，暂时不处理读写过程，所以当被动建立连接的一方处于 `SYN_RECV` 状态时，收到来自主动建立连接一方的 `ACK` 信号后，会唤醒 `wait_accept` 上的子套接字。

关闭过程

```
1 void tcp_sock_close(struct tcp_sock *tsk)
2 {
3     // fprintf(stdout, "TODO: implement %s please.\n", __FUNCTION__);
4     switch(tsk->state) {
5         case TCP_CLOSED:
6             break;
7         case TCP_LISTEN:
8             tcp_unhash(tsk);
9             tcp_bind_unhash(tsk);
10            tcp_set_state(tsk, TCP_CLOSED);
11            break;
12        case TCP_SYN_SENT:
13            tcp_unhash(tsk);
14            tcp_set_state(tsk, TCP_CLOSED);
15            break;
```

```

16     case TCP_ESTABLISHED:
17         tcp_set_state(tsk, TCP_FIN_WAIT_1);
18         tcp_send_control_packet(tsk, TCP_FIN);
19         break;
20     case TCP_CLOSE_WAIT:
21         tcp_set_state(tsk, TCP_LAST_ACK);
22         tcp_send_control_packet(tsk, TCP_FIN);
23         break;
24     default:
25         break; // temp deal
26 }
27 }

```

要考虑在任何场景下某一方选择 close 整个套接字，因为这是 TCP 应用用来彻底释放资源的唯一调用。由于本例比较简单，而且数据传输不会出错，所以不用考虑得太多，关注资源的释放即可。

实验结果与分析

我在实验代码中加入了一些 DEBUG 信息用来检测套接字的状态。利用实验给定的 server 和 client 小应用，可以看到两边套接字的状态切换过程，从而验证 TCP 实现的正确性：

```

1  运行给定网络拓扑(tcp_topo.py)
2  在节点h1上执行TCP程序
3  执行脚本(disable_tcp_rst.sh, disable_offloading.sh)，禁止协议栈的相应功能
4  在h1上运行TCP协议栈的服务器模式  (./tcp_stack server 10001)
5  在节点h2上执行TCP程序
6  执行脚本(disable_tcp_rst.sh, disable_offloading.sh)，禁止协议栈的相应功能
7  在h2上运行TCP协议栈的客户端模式，连接至h1，显示建立连接成功后自动关闭连接 (./tcp_stack
    client 10.0.0.1 10001)

```

结果：

```

root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/14-tcp_stack# ./tcp_stack server 100
01
DEBUG: find the following interfaces: h1-eth0.
Routing table of 1 entries has been loaded.
DEBUG: 0.0.0.0:10001 switch state, from CLOSED to LISTEN.
DEBUG: listen to port 10001.
DEBUG: 0.0.0.0:10001 current state is LISTEN.
DEBUG: 10.0.0.1:10001 switch state, from CLOSED to SYN_RECV.
DEBUG: 10.0.0.1:10001 current state is SYN_RECV.
DEBUG: 10.0.0.1:10001 switch state, from SYN_RECV to ESTABLISHED.
DEBUG: accept a connection.
DEBUG: 10.0.0.1:10001 current state is ESTABLISHED.
DEBUG: 10.0.0.1:10001 switch state, from ESTABLISHED to CLOSE_WAIT.
DEBUG: 10.0.0.1:10001 switch state, from CLOSE_WAIT to LAST_ACK.
DEBUG: 10.0.0.1:10001 current state is LAST_ACK.
DEBUG: 10.0.0.1:10001 switch state, from LAST ACK to CLOSED.

```

```
root@ubuntu:/mnt/hgfs/[...]/[...]/[...]/14/14-tcp_stack# ./tcp_stack client 10.0.0.1 10001
DEBUG: find the following interfaces: h2-eth0.
Routing table of 1 entries has been loaded.
DEBUG: 10.0.0.2:12345 switch state, from CLOSED to SYN_SENT.
DEBUG: 10.0.0.2:12345 current state is SYN_SENT.
DEBUG: 10.0.0.2:12345 switch state, from SYN_SENT to ESTABLISHED.
DEBUG: 10.0.0.2:12345 switch state, from ESTABLISHED to FIN_WAIT-1.
DEBUG: 10.0.0.2:12345 current state is FIN_WAIT-1.
DEBUG: 10.0.0.2:12345 switch state, from FIN_WAIT-1 to FIN_WAIT-2.
DEBUG: 10.0.0.2:12345 current state is FIN_WAIT-2.
DEBUG: 10.0.0.2:12345 switch state, from FIN_WAIT-2 to TIME_WAIT.
ERROR: invalid tcp packet while in state FIN_WAIT_2
DEBUG: 10.0.0.2:12345 switch state, from TIME_WAIT to CLOSED.
```